

InventOR: A Human-Reviewed, Prompt-Configured Large Language Model Decision Support System for Material-Level Inventory Control

Aris Dressino^{a,*}

^a*TUM School of Management, Technical University of Munich, Arcisstraße
21, Munich, 80333, Germany*

Abstract

Material planners must set a reorder point, a safety stock and an order quantity for thousands of plant-material pairs, usually under ERP parameters that were last revised long ago. We define and evaluate *InventOR*, a decision support system (DSS) for this recurring decision. For each material, a large language model (LLM) configured only through a prompt, without training or fine-tuning, proposes a reorder-point/order-quantity (r, Q) policy with a written justification and an escalation flag. Deterministic software gates the proposal against minimum order quantities and storage limits and replays it on held-out demand before the planner, who remains the decision-maker, reviews it. We derive six design requirements and evaluate the artifact on 365 plant-material pairs from three plants of an industrial manufacturer, using three independent LLM runs, seven reference arms and an ablation that withholds the ERP planning parameters. In a hard-limit storage replay, the LLM policies cost about 11% less than ERP-derived parameters at a mean fill rate 0.1–0.4 points lower, but 6%, 8% and 18% more than

*Corresponding author. Telephone: +49 89 289 25000. ORCID: 0009-0003-0693-9051.
Email address: aris.dressino@tum.de (Aris Dressino)

history-only, Croston/SBA and empirical-quantile (r, Q) baselines. The saving comes from carrying 13–17% less inventory with 11–20% more stockout days; when shortages are priced at one unit cost, ERP-derived parameters are cheaper and a simulation-tuned policy is cheapest. Static storage checks are insufficient: simulated trajectories exceed storage limits for 53–66% of pairs. InventOR is a checkable, human-reviewed candidate generator, not an optimizer that beats data-driven baselines; benefits to planners are untested. We contribute design principles, review rules and a research agenda.

Keywords: Decision Support Systems, Large Language Models, Generative Artificial Intelligence, Inventory Control, Human–AI Decision Making, Design Science, Operations Research

1. Introduction

Replenishment parameters in manufacturing are set per plant-material pair: a reorder point, a safety stock and an order quantity. In large ERP systems these values often remain at levels set years earlier, while demand, supplier performance and storage constraints change. Revisiting them requires analysts who can fit demand and lead-time statistics, apply inventory theory, respect minimum order quantities (MOQs) and storage limits, and explain the result to planners. A conversational large language model (LLM) with a code interpreter appears to offer this scarce skill, and a growing literature explores such tools for supply-chain and operations decisions [1, 2, 3, 4, 5].

Three problems limit its use. First, LLM outputs can be unfaithful or fabricated [6], so a proposed policy is a candidate, not a decision. Second, superiority claims over ERP parameters are hard to interpret when the LLM and the comparators receive different information. Third, a storage rule that

holds when a policy is written down can be violated as stock builds up over time.

This paper treats the problem as one of decision support system (DSS) design. We define *InventOR* as a human-reviewed, model-driven DSS for the recurring decision of setting replenishment parameters per plant-material pair (Section 2). A planner remains the decision-maker. The system supplies, for each material, a candidate (r, Q) policy proposed by an enterprise LLM application that is configured only through a prompt and receives no project-specific training or fine-tuning, together with a justification, an escalation flag and a counterfactual simulation; deterministic software parses, gates and simulates every proposal before the planner reviews it and alone decides whether any ERP parameter is changed.

We follow a design-science logic [7, 8]: we state problem and scope, derive design requirements, build the artifact and evaluate it on real industrial data. The questions are how an LLM should be embedded so that its proposals are bounded, verifiable, reproducible and reviewable; what cost and service they achieve when the information available to each approach is controlled; whether they respect storage limits along simulated trajectories and are stable across runs; and when a planner should trust, escalate or reject a proposal.

Our contributions are: (C1) a definition and scope of the DSS with six design requirements that bound the role of the LLM (Section 2); (C2) an architecture and operating rules that separate generation, deterministic gating, simulation and planner review, with design principles (Sections 4 and 8); (C3) an information-controlled, trajectory-level evaluation with seven reference arms, an ablation denied the ERP planning parameters and a hard-limit storage replay (Section 5); and (C4) complete-cohort evidence from

three LLM runs on 365 plant-material pairs, with managerial and research implications.

Our main findings are as follows. Against ERP-derived parameters the LLM policies are about 11% cheaper at a mean fill rate 0.1–0.4 points lower, but the LLM costs 6%, 8% and 18% more than history-only, Croston/SBA and empirical-quantile (r, Q) baselines at essentially equal service, and its saving over ERP comes from carrying less inventory at the price of more stockout days. Storage limits that hold in a static check are violated along unconstrained trajectories for 53–66% of pairs. The simulation therefore does not show superior optimization by the LLM; it shows that a prompt-configured LLM can produce checkable candidate policies without code or training, whose benefit to planners remains to be tested.

2. Problem, Scope and Definition

2.1. The decision problem

The decision supported is the periodic (re)setting of a reorder point, a safety stock and an order quantity for each plant-material pair held in an ERP system. In the classification of Gorry and Scott Morton [9] it is a partly structured operational-control decision: the calculation is textbook, but the choice of estimator, the treatment of intermittent or trending demand, the plausibility of recorded lead times and the handling of constraints require judgement that is applied inconsistently when thousands of materials compete for a few analysts' time. The decision-maker is the material planner; hard constraints are the minimum order quantity (MOQ) and a physical storage limit; and feedback is slow, because a changed parameter shows its effect only in later demand. The scope is one pair at a time (single item, single

echelon, continuous review, lost sales, (r, Q) family); order release, supplier selection, forecasting-model development and multi-echelon networks are out of scope, and the LLM is a candidate generator that does not forecast, place orders, simulate or decide.

2.2. Definition and level of automation

Definition. *InventOR is a human-reviewed, model-driven decision support system that, for one plant-material pair at a time, (i) obtains from a large language model, configured only by a written prompt, a candidate (r, Q) policy with a justification and an escalation flag; (ii) verifies the candidate deterministically against explicit constraints; (iii) evaluates it by counterfactual simulation against reference policies on held-out demand; and (iv) presents the result to a planner, who alone decides whether ERP parameters are changed.*

In the classical DSS architecture of data, model and dialog components [10, 11], the LLM and the simulator form the model component, the material-scoped data extract and the prompt form the data and knowledge components, and the review interface is the dialog component (Section 4.1). In terms of the levels of automation of Parasuraman et al. [12], InventOR operates at a low-to-intermediate level: it selects and suggests one alternative and leaves execution to the human. Accepting automated suggestions invites complacency and automation bias [13], which the review rules of Section 4.3 are meant to counter.

90 2.3. Design requirements

91 The requirements follow from the known unreliability of LLM outputs
 92 [6], guidance on human–AI interaction and accountability [14, 15], run-to-
 93 run nondeterminism of LLM inference [16] and the constraints of inventory
 94 control. Table 1 ties each to a design feature (Section 4) and an evaluation
 95 (Section 6).

Table 1: Design requirements, the design features that address them and where they are evaluated.

ID	Requirement	Design feature	Evaluation
DR1	Bounded generation: the LLM proposes and has no authority	Narrow candidate-generator role; prohibited methods; no order placement	Sections 6.1, 6.3
DR2	Verifiable feasibility along time, not only at a point	Deterministic gate; hard-limit trajectory replay	Sections 6.1, 6.5
DR3	Counterfactual evidence for the reviewer	Held-out simulation against reference arms	Sections 6.2, 6.8
DR4	Information transparency: comparisons control what each arm sees	Seven reference arms; LLM ablation without ERP parameters	Sections 6.3, 6.4
DR5	Reproducibility and accountability of the configuration	Frozen prompt file; per-call log; repeated runs	Section 6.7
DR6	Human control with explicit escalation	Trust, escalate and reject rules; escalation flag	Sections 4.3, 6.6

96 3. Related Work and Positioning

97 *Decision support systems and human–AI decision making.* DSS research
 98 has long distinguished data-, model- and knowledge-driven systems and has
 99 stressed that the value of a DSS lies in improving the judgement of a hu-
 100 man decision-maker rather than in replacing it [9, 10, 11]. Research on
 101 human–automation interaction shows that the way suggestions are scoped,
 102 presented and verified determines whether users rely on them appropriately
 103 or over-rely on them [12, 13],. Recent DSS work derives design principles
 104 for LLM-based assistants [17], reports design and pilot implementations of
 105 AI-enabled decision support in operations [18], and shows that the uptake of

106 algorithmic advice depends on transparency and on the user’s metacognition
107 [19, 20]; evaluations of data-driven forecasting systems in practice likewise
108 stress lessons beyond accuracy [21]. Model-driven support also shapes how
109 groups handle disagreement [22], and human override of algorithmic safety-
110 stock advice is a documented risk [23]; InventOR therefore keeps the planner
111 as decision-maker and makes the evidence for each proposal explicit.

112 *LLMs in operations and supply-chain decisions..* Reviews and frameworks
113 map the opportunities for generative AI in supply chains [1, 2, 24, 25, 3, 4, 26]
114 and stress that adoption depends on assigned human roles, data quality and
115 governance [5, 27, 28, 29, 30]. Evidence from manufacturing planning is thin:
116 a fine-tuned LLM for digital-manufacturing optimization [31], a transformer
117 trained to set order timing and quantity and benchmarked against (r, Q) [32],
118 and experiments showing that ChatGPT reproduces human decision biases
119 of managers [33]. Most technical LLM work translates problem statements
120 into optimization models and solver code [34, 35, 36, 37], studies multi-agent
121 control in inventory flows [38, 39], disruption modelling [40] or classical OR
122 tasks such as routing [41]. The closest LLM study to ours uses an LLM
123 as a forecasting copilot for spare-parts demand, combined with simulation,
124 and reports high variability for irregular demand that can be reduced by
125 averaging repeated queries and enriching prompts [42]. We test run-to-run
126 variability directly (Section 6.7); the difference is that our LLM proposes a
127 complete policy, not a forecast.

128 *Data-driven inventory control without LLMs..* Learning-based replen-
129 ishment is well established, including data-driven policies learned from
130 sequentially observed non-stationary demand [43] and forecasts augmented
131 with leading indicators [44], and heuristics remain competitive in constrained

132 settings such as perishable inventory [45]. Deep reinforcement learning has
 133 been applied to omnichannel and multi-echelon systems [46, 47, 48] and
 134 aerospace replenishment [49]; deep-learning forecasting has been integrated
 135 with replenishment [50]; contextual distributionally robust optimization
 136 uses side information to set order quantities [51]; and distribution-free
 137 online algorithms handle capacitated warehouses under uncertain price and
 138 demand [52]. The closest operational neighbour is a trained data-driven
 139 framework that jointly sets replenishment parameters, safety stock and
 140 safety lead time on real automotive data [53]. Reinforcement learning has
 141 also been combined with demand forecasting to optimize order points [54].
 142 These methods need training or tuning; InventOR needs none, but its
 143 generator cannot be inspected beyond its written justification.

144 *Intermittent and spare-parts demand..* Material histories in manufacturing
 145 are frequently intermittent or lumpy. The statistical literature offers special-
 146 ized estimators and classification schemes [55, 56, 57, 58, 59], and recent stud-
 147 ies develop feature-based forecast combinations and temporal aggregation for
 148 intermittent demand [60, 61], forecasting for spare parts across product life-
 149 cycle phases [62], and integrated forecasting and planning or multi-echelon
 150 stocking models for repairable and spare parts [63, 64].

151 InventOR is not a new inventory model and does not claim to beat
 152 trained or optimized methods; its (r, Q) policy is the standard workhorse
 153 of continuous-review control [65, 66, 67].

154 4. The InventOR Decision Support System

155 4.1. Architecture and components

156 Following the classical DSS components [10, 11], InventOR has a *data*
157 component (a material-scoped extract of daily demand, lead times, working-
158 day flags, cost and constraint parameters and ERP reference values; the
159 fields supplied are controlled and logged; DR4), a *knowledge* component (the
160 system prompt file defining the task, the prohibited methods, the storage
161 rules and the output schema, Supplementary File S1; DR1, DR5), a *model*
162 component split into a generator (an enterprise LLM application with a code
163 interpreter, configured only by the prompt, proposing one (r, Q) policy with
164 a justification and an escalation flag; DR1) and a deterministic verifier and
165 simulator (parser, gate, held-out simulation with hard storage-limit replay
166 and reference arms; DR2, DR3), a *dialog* component (policy, justification,
167 flag and simulated comparison shown to the planner under the trust, escalate
168 and reject rules; DR6) and a *governance* layer (frozen prompt file, per-call
169 log of input hash, time, session and raw response, repeated runs; DR5). The
170 generator is the only non-inspectable component. In the evaluated prototype
171 the dialog component is the per-pair output record; no interactive interface
172 was built or evaluated.


173 4.2. Operation in four steps

174 For one material, InventOR answers how much safety stock to hold, at
175 what level to reorder and how much to order, in four steps (Figure 1).

176 **Step 1. Generate.** The LLM, accessed through an enterprise application
177 with a code interpreter, receives one material’s history (daily de-
178 mand, realized lead times, working-day flags) and a short parameter

InventOR prototype workflow

Executed material-level path from prepared history to planner-reviewable simulation output.



Deterministic components reproduce simulation results from stored artifacts; LLM inference itself may vary.

Figure 1: InventOR workflow: material data are prepared, an LLM proposes a policy, deterministic software parses, checks and simulates it, and a planner reviews the result.

block (unit cost, holding rate, ordering cost, MOQ, service target, storage limit, and current ERP safety stock and safety lead time for comparison). A written prompt (the only configuration; no weights are changed) [68] instructs it to compute demand and lead-time statistics, apply an economic-order-quantity (EOQ) argument and a normal-approximation safety stock, respect the storage limit, and report a structured policy with a short justification and escalation flag.

Step 2. Gate. Deterministic code extracts the policy, rejects malformed or infeasible outputs (missing or non-numeric fields, non-positive reorder point or order quantity, negative safety stock, reorder point below safety stock, safety stock plus order quantity above the storage limit) and retries the generation up to three times. The backtest applies one further admission check, that the order quantity is not below the MOQ.

Step 3. Simulate. Each accepted policy is replayed on the material’s demand observed *after* the data cutoff, a held-out evaluation in the

196 sense of out-of-sample forecast testing [69, 70], together with the
 197 comparator policies, producing holding and ordering cost, fill rate
 198 and stockout days.

199 **Step 4. Review.** A planner sees the policy, the LLM’s justification and
 200 flag, and the simulated comparison. Only the planner changes ERP
 201 parameters.

202 The LLM is a *candidate generator* whose output is never trusted without
 203 the gate; it does not forecast, place orders or run the simulation. The prompt
 204 prohibits Monte Carlo simulation, distribution fitting, machine-learning li-
 205 braries and specialized intermittent-demand estimators, so that the policy
 206 rests on transparent statistics.

207 4.3. *Trust, review and reject rules*

208 We propose three rules for planners, in line with guidelines for human–AI
 209 interaction and AI auditing [14, 15]. The rules and the 0.5-point threshold
 210 are design proposals not evaluated with planners; the trust rule would be
 211 met by a minority of pairs (the LLM is cheaper than the history-only OR in
 212 74–77 of 346).

213 **Trust (apply after review)** when the proposal passes the gate, the sim-
 214 ulated trajectory respects the storage limit when orders are truncated at it,
 215 and the simulated cost is not above that of the best history-only OR ref-
 216 erence (Croston/SBA or empirical quantile) with a fill rate not more than
 217 about 0.5 points lower.

218 **Escalate** when the LLM raises an escalation flag, the material is lumpy or
 219 has strongly seasonal or trending demand, the plant or calendar data were
 220 repaired, or the policy differs strongly between repeated generations.

221 **Reject** when the proposal fails the gate after the retry budget, exceeds the
222 storage limit in the hard-limit replay, or is dominated (higher cost and lower
223 fill) by a history-only reference.

224 *4.4. Provenance and reproducibility of the configuration*

225 The system prompt, a plain-text file deployed unchanged in the enter-
226 prise application, defines the task, forbidden methods, capacity rules and
227 output schema (Supplementary File S1). Every call logs the SHA-256 of its
228 input, UTC time, session identifier and raw response. The platform does
229 not expose the model identifier or sampling settings; the evaluated model
230 (Anthropic Claude Sonnet 4.6 with a code interpreter) is operator-declared.
231 Reproducibility rests on the prompt file, the call logs and three independent
232 runs (Section 6.7).

233 **5. Application and Evaluation Design**

234 *5.1. Evaluation approach*

235 The evaluation addresses the design requirements (Table 1). The pri-
236 mary basis is the hard-limit replay with working-day receipts and no short-
237 age penalty; sensitivity analyses vary these choices. Following design science
238 [7, 8], the evaluation is artificial and ex post in the framework of Venable
239 et al. [71]: real data and a faithful simulation but no live users. It estab-
240 lishes technical efficacy, feasibility and reliability, not usability, adoption or
241 improved planner decisions.

242 *5.2. Industrial application and data*

243 The application is a pilot at a large manufacturer. The dataset contains
244 233,078 daily records for 411 plant-material pairs from three plants (A, B and

245 C), with a data cutoff on 10 March 2026 and a post-cutoff horizon observed
 246 to 27 August 2026. Of the 411 pairs, 365 are eligible for scoring (234 in
 247 Plant A, 65 in Plant B, 66 in Plant C) and are sent to the LLM. For 19
 248 pairs (all in Plant A) the MOQ plus the reference safety stock exceeds the
 249 storage limit, so no comparator can satisfy both rules; these are excluded
 250 from outcome comparison, leaving a *common-feasible cohort* of 346 pairs
 251 (215, 65 and 66). Demand is concentrated (the five largest materials account
 252 for 40.3% of demand). The source working-day flag of Plant C is all zero; we
 253 reconstruct it from the same-date majority of Plants A and B and test this
 254 repair in Section 6.8. Data quality is a known threat to decision support [72]
 255 and inventory records [73, 74, 75].

256 Each pair is sent as a material-scoped CSV with daily demand, recorded
 257 lead times, a working-day flag, the minimum order quantity, unit cost,
 258 holding-cost rate, ordering cost, service-level target, the storage limit, and
 259 two *ERP planning parameters* (current safety stock and safety lead time),
 260 which are withheld in the ablation arm. The CSV also carries supplier and
 261 lane context fields that the prompt does not use.

262 5.3. Reference arms and information content

263 ERP safety stock and safety lead time embed planner judgement and are
 264 alternative ways of buffering unreliability [76], and planned lead times can
 265 drift upward [77]. The LLM and the classical comparators therefore do not
 266 necessarily use the same information, and we construct arms that make it
 267 explicit (Table 2).

268 Three further arms use no ERP input and represent stronger data-driven
 269 practice. The *Croston/SBA OR* replaces the mean daily demand by the
 270 Croston estimate with the Syntetos–Boylan bias correction [55, 56] (smooth-

Table 2: Comparison arms and the information each uses. All arms respect MOQ and are projected to the storage limit when possible.

Arm	Parameters from	ERP used	inputs	Purpose
SAP-derived	ERP safety stock; EOQ order quantity	Safety stock		ERP-derived reference (not observed practice)
SLT-informed OR	Stationary (r, Q) with ERP safety lead time added to the lead time	Safety lead time		OR policy using the ERP safety lead time
ERP-floor OR	SLT-informed OR with safety stock floored at the ERP value	Both		OR arm matched to the LLM’s full information
History-only OR	Stationary (r, Q) from demand and lead-time history, mean lead time	None		ERP-independent baseline
Croston/SBA OR	(r, Q) with the Syntetos–Boylan-approximated Croston demand rate and its forecast-error spread, mean lead time	None		Intermittent-demand baseline
Empirical-quantile OR	(r, Q) with the reorder point set to the service-level quantile of historical rolling lead-time demand	None		Distribution-free baseline
Simulation-tuned (r, Q)	Reserve factor and order-quantity multiple chosen by simulation on the last 250 working days before the cutoff	None		Trained data-driven baseline
LLM (runs 1–3)	LLM proposal with all CSV fields (Section 5.2)	Both		Main treatment
LLM without ERP	LLM with ERP safety stock and safety lead time set to “not supplied”	None		Ablation: what the LLM does when it cannot anchor on ERP

271 ing constant 0.1), uses the root mean squared one-step forecast error as
 272 demand spread and the history-only lead-time moments, which suits inter-
 273 mittent demand [57]. The *empirical-quantile OR* sets the reorder point to the
 274 service-level quantile of historical rolling lead-time demand over the rounded
 275 mean lead time. The *simulation-tuned (r, Q)* policy selects, from a grid of
 276 reserve factors and order-quantity multiples around the history-only policy,
 277 the cheapest candidate whose fill rate on the last 250 working days *before*
 278 *the cutoff* meets the service target (or the best-fill candidate if none does).

279 No post-cutoff data are used in any of these arms. We do not implement a
 280 deep reinforcement-learning baseline [47, 46]; this is a limitation (Section 8).

281 5.4. *Simulation, capacity replay and metrics*

282 Each pair is simulated independently on the post-cutoff horizon as a
 283 single-item, single-echelon continuous-review (r, Q) system [65, 78, 79, 67, 66]
 284 with lost sales [80], in the spirit of simulation-based experiments on supply
 285 and timing uncertainty [81], a shared opening stock, an empty pipeline and
 286 arrivals after the rounded mean lead time in working days (Appendix Ap-
 287 pendix A). MOQs are treated as hard order-size constraints, as imposed
 288 by suppliers [82], and storage limits as capacity bounds on the replenish-
 289 ment decision [52]. The primary metrics are the mean material-horizon
 290 holding-plus-ordering cost, the mean material-level fill rate, the aggregate
 291 (volume-weighted) fill rate and the number of stockout days. We report
 292 paired differences with 10,000-sample bootstrap 95% intervals.

293 A static check $SS + Q \leq \text{max_storage_units}$ does not show that storage
 294 limits hold as inventory evolves, because the reorder point includes lead-time
 295 demand and inventory on order can exceed the limit on arrival. We therefore
 296 record the post-receipt on-hand stock for every policy and day and report
 297 the share of pairs and days above the limit (the *unconstrained* simulation),
 298 and we *replay* every policy with each order truncated to the storage limit
 299 minus the current inventory position (the *hard-limit replay*), which is our
 300 primary evaluation.

301 5.5. *Repeated runs and sensitivities*

302 To examine run-to-run variation, the entire cohort is generated three
 303 times (runs 1, 2 and 3) with the same prompt file and the same data. Run 1

304 serves as the reference for tabulated results unless stated otherwise; runs 2
305 and 3 provide replicates. Sensitivity analyses vary the working-day recon-
306 struction (majority, any-plant, strict-all, exclusion of Plant C and calendar-
307 day arrivals), the shortage penalty (zero, one and five times unit cost).

308 6. Results

309 6.1. Completion and policy characteristics

310 All four runs (three with the full input and one ablation, Section 6.4)
311 return a valid, storage-feasible (r, Q) policy for all 365 pairs after at most
312 three attempts per pair. The validity gate rejects non-positive reorder points
313 and order quantities, negative safety stocks and reorder points below the
314 safety stock. The ERP safety stock is zero for 198 of the 365 pairs (188 of
315 the 346 in the cohort below), so for a majority of pairs the SAP-derived arm
316 holds no safety stock at all.

317 6.2. Headline comparison with ERP-derived parameters

318 Unless stated otherwise, all results use the hard-limit replay on the
319 346-pair common-feasible cohort with working-day receipts and no shortage
320 penalty (Section 6.5 gives the unconstrained simulation). Table 3 compares
321 all arms; its last columns give the LLM cost relative to each arm in the three
322 runs (positive: the LLM costs more). In run 1 the LLM policies cost 11.4%
323 less than SAP-derived policies and 30.5% less than the SLT-informed OR
324 (mean 2,928 versus 3,305 and 4,211), at a mean fill 0.18 points lower than
325 SAP-derived (98.11% versus 98.30%; interval -0.54 to $+0.10$). The mean
326 cost difference against SAP-derived is -377 (95% interval -651 to -154);
327 the LLM is cheaper in 241 of 346 pairs and, in 226, simultaneously no more
328 expensive and no worse in fill. It incurs more stockout days (910 versus 804).

Relative to SAP-derived parameters the LLM carries about 17% less average inventory (16,256 versus 19,565 units) and about 16% lower holding cost at about the same number of orders per year (17.3 versus 17.6), and it accepts 11–20% more stockout days (runs 3 and 2). The saving is therefore a trade of inventory for service, not a better ordering rhythm, and Section 6.8 shows that once shortages carry a cost this trade no longer favours the LLM.

Table 3: Hard-limit replay, common-feasible cohort ($n = 346$), working-day receipts, zero shortage penalty. Cost is the mean material-horizon holding-plus-ordering cost; average inventory is in units; orders are per year. The last column gives the LLM cost relative to the arm in runs 1/2/3 (positive: the LLM costs more).

Arm	Mean cost	Holding	Ordering	Avg. inv.	Orders	Mean fill (%)	Aggr. fill (%)	Stockout days	LLM vs arm (%)
ERP-derived (SAP)	3,305	2,172	1,132	19,565	17.6	98.30	79.34	804	−11.4 / −11.3 / −11.4
SLT-informed OR	4,211	2,608	1,603	21,614	24.5	98.23	79.15	831	−30.5 / −30.4 / −30.5
ERP-floor OR	4,248	2,633	1,616	21,937	24.7	98.23	79.15	831	−31.1 / −31.0 / −31.1
History-only OR	2,764	1,725	1,039	15,972	16.3	98.14	78.72	874	+5.9 / +6.0 / +5.9
Croston/SBA OR	2,705	1,665	1,040	15,521	16.3	98.30	78.87	803	+8.2 / +8.4 / +8.3
Empirical-quantile OR	2,489	1,543	946	13,331	14.8	98.35	78.87	747	+17.7 / +17.8 / +17.7
Simulation-tuned (r, Q)	3,266	2,473	793	11,914	12.6	96.05	84.84	1,973	−10.4 / −10.3 / −10.3 [†]
LLM (run 1)	2,928	1,824	1,104	16,256	17.3	98.11	78.87	910	
LLM (run 2)	2,931	1,850	1,081	16,599	17.0	97.89	78.70	968	
LLM (run 3)	2,929	1,841	1,088	16,958	17.0	98.16	78.90	890	
LLM without ERP	3,014	1,902	1,112	16,507	17.3	98.06	78.89	890	−8.8 [†] (vs SAP)

[†] Other comparisons of the ablation are in Section 6.4. [‡] Different operating point (mean fill 96.05%, aggregate fill 84.8%); the paired interval of the cost difference includes zero, and with priced shortages this policy is the cheapest arm (Section 6.8).




Figure 2: Cost structure and effect of pricing shortages (hard-limit replay, 346 pairs; LLM rows in bold). (A) Mean cost per pair split into holding and ordering cost, with the total at the end of each bar. The LLM and the data-driven ORs carry less inventory than SAP-derived parameters and the ERP-informed ORs. (B) Total cost when lost sales are charged at one and five times the unit cost, relative to SAP-derived parameters (negative: cheaper). Once shortages are priced SAP-derived parameters beat the LLM, and the simulation-tuned policy is cheapest.

335 *6.3. Information asymmetry: what the LLM adds beyond the ERP inputs*

336 First, the SLT-informed OR and its ERP-floor variant are the most ex-
337 pensive arms (4,211 and 4,248): the ERP safety lead time inflates their
338 effective lead time and buffers. Against them the LLM is 30–31% cheaper in
339 every run, which mostly measures how heavily ERP inputs over-buffer.

340 Second, the *history-only OR*, which uses no ERP inputs, costs 2,764 at
341 a fill of 98.14%. The LLM costs 5.9–6.0% more in all three runs (run 1:
342 +164, interval +91 to +240), with fill differences of -0.03 , -0.25 and $+0.01$
343 points whose intervals include zero, and is cheaper in only 74–77 of 346 pairs.
344 The stronger data-driven arms widen the gap: the LLM costs 8.2–8.4% more
345 than the Croston/SBA OR (2,705, fill 98.30%; run 1: +223, interval +97
346 to +362; cheaper in 91–99 pairs) and 17.7–17.8% more than the empirical-
347 quantile OR (2,489, fill 98.35%; +439, interval +329 to +557; cheaper in
348 44–49 pairs; mean fill 0.24 points lower). Part of the LLM’s advantage over
349 ERP-derived parameters therefore reflects weak ERP benchmarks.

350 The simulation-tuned (r, Q) policy sits at a different operating point: an
351 aggregate fill of 84.8% and the lowest inventory (11,914 units) but a mean
352 fill of 96.05% and 1,973 stockout days, because it is tuned on the pre-cutoff
353 year. The LLM costs 10.3–10.4% less at zero penalty with a 2.1-point higher
354 mean fill, but the interval includes zero ($-2,527$ to $+857$); once shortages
355 are priced the tuned policy is by far the cheapest arm (Section 6.8).

356 Third, the LLM improves on ERP-derived parameters (11.3–11.4%
357 cheaper at 0.14–0.40 points lower fill) but not on the best classical baselines,
358 which also yield a policy for every material at lower cost.

359 6.4. LLM without ERP planning parameters (ablation)

360 In this ablation the LLM receives the same inputs as in runs 1–3, but the
361 ERP safety stock and safety lead time are blanked in the CSV and marked
362 “not supplied”; the system prompt file is unchanged. It tests whether the
363 LLM anchors on the ERP planning parameters. All 365 pairs returned a
364 valid, storage-compliant policy, and the policies change only slightly beyond
365 run-to-run variation: median ratios to run 1 are 1.00, and the reorder point
366 differs by more than 10% in 54.6% of pairs (48.3% and 48.0% between ERP-
367 supplied runs). In the hard-limit replay the cost is 3,014 at a mean fill of
368 98.06%, against 2,928–2,931 in runs 1–3. It is 8.8% below SAP-derived pa-
369 rameters (interval -566 to -61 ; fill -0.24 points, interval including zero) and
370 28.4–29.1% below the SLT-informed and ERP-floor ORs, but 9.0% above the
371 history-only OR (interval $+134$ to $+387$; cheaper in 91 of 346 pairs), 11.4%
372 above the Croston/SBA OR, 21.1% above the empirical-quantile OR and
373 7.7% below the simulation-tuned policy. By plant the pattern is unchanged
374 (-10.8% in Plant A, -21.9% in Plant B, $+41.6\%$ in Plant C against $+36.2\%$
375 in run 1), so the Plant C loss does not stem from reliance on ERP parameters.
376 The information-asymmetry concern therefore does not explain the LLM’s
377 advantage over ERP-derived parameters, and the ERP planning fields ap-
378 pear to be of limited use to the LLM. The ablation withholds only these
379 two fields; unit cost, holding rate, ordering cost, MOQ, service target and
380 storage limit remain available, as they are for all reference arms.

381 6.5. Storage limits along simulated trajectories

382 All 365 LLM policies satisfy the static rule $SS+Q \leq \text{max_storage_units}$,
383 yet in the unconstrained simulation post-receipt on-hand stock exceeds the
384 limit for 183 pairs (52.9%; 178–183 across runs; 10.2% of pair-days) for the

385 LLM, 198 (57.2%) for SAP-derived and 230 (66.5%) for the SLT-informed
 386 OR, because inventory on order and lead-time demand add to stock above
 387 $SS + Q$. The static check is thus necessary but not sufficient for any
 388 method. Under the hard-limit replay no arm exceeds the limit by more
 389 than one unit; enforcing it raises the LLM mean cost from 2,850 to 2,928
 390 and SAP-derived from 3,184 to 3,305 without changing the ranking.

391 *6.6. Plants and demand regimes*

392 The average hides a strong plant effect (Figure 3, hard-limit replay,
 393 run 1). The LLM costs 14.1% less than SAP-derived in Plant A (cheaper in
 394 179 of 215 pairs) and 22.7% less in Plant B (54 of 65), but 36.2% *more* in
 395 Plant C (cheaper in 8 of 66); the Plant C loss is similar in the other full runs
 396 (+32.2%, +31.5%) and persists under every working-day reconstruction in
 397 which Plant C can be scored (Section 6.8). Against the data-driven base-
 398 lines the LLM costs more in every plant: 7.2%, 3.6% and 6.7% more than the
 399 history-only OR, 13.1%, 2.3% and 6.2% more than the Croston/SBA OR,
 400 and 20.8%, 7.7% and 29.5% more than the empirical-quantile OR in Plants
 401 A, B and C. The SLT-informed and history-only ORs coincide in Plant C
 402 because the ERP safety lead time is zero for all its pairs.

403 By demand regime, the 33 lumpy materials show the largest reduction
 404 against SAP-derived (−25.4% in run 1) and the 313 smooth materials −9.8%
 405 at a 0.24-point lower mean fill. On lumpy materials the service effect is
 406 inconsistent (mean fill +0.38 points in run 1, −0.29 in run 2), the history-only
 407 and Croston/SBA ORs cost about 9–10% less than the LLM, and only the
 408 empirical-quantile OR is slightly more expensive (1.2%); with 33 materials
 409 this is indicative only. By terciles of the coefficient of variation, the LLM
 410 costs 15.7%, 7.5% and 12.9% less than SAP-derived (lowest to highest), with

411 a 0.68-point lower mean fill in the middle tercile.




Figure 3: Where the LLM wins and loses (run 1, hard-limit replay). Each marker is the LLM’s mean cost relative to one comparator in a segment of pairs (number of pairs in brackets); markers to the right of the vertical line mean the LLM costs more. The LLM is cheaper than SAP-derived parameters in Plants A and B but not in Plant C, and costs more than the three data-driven baselines in almost every segment.

412 6.7. Run-to-run stability

413 Across the three runs the aggregate results are very similar: the LLM
414 cost relative to SAP-derived is -11.4% , -11.3% and -11.4% , and the policy-
415 value medians are close (reorder point 1,584, 1,582 and 1,466; order quantity
416 2,737, 2,777 and 2,693; safety stock 683, 688 and 669). Individual policies
417 vary, however. Between runs 2 and 3 on the 365 pairs, all three values are
418 identical for only 36 pairs; every policy resolves to the (r, Q) family, and the
419 median absolute differences are 81 units (reorder point), 6 (order quantity)
420 and 3 (safety stock). The reorder point differs by more than 10% in about
421 48% of pairs, and a few pairs vary strongly (one reorder point was 14,882,
422 14,947 and 29,218 units). Aggregate behaviour is reproducible, but a single
423 generation per material is a noisy estimate that should be reviewed [16].

424 6.8. Sensitivity analyses

425 *Plant C calendar..* Plant C’s working-day flag required reconstruction, so we
426 re-scored all policies under four alternatives (any-plant, strict-all, exclusion
427 of Plant C, calendar-day arrivals). Where Plant C can be scored, the LLM
428 remains 31–35% more expensive than SAP-derived there in the three full
429 runs, and the cohort-wide advantage is -10.8% to -11.1% ; the strict and
430 exclusion variants (280 pairs) show -17.0% to -17.6% . The Plant C result
431 is therefore not an artefact of the repair, although the LLM saw demand
432 series built on the repaired calendar.

433 *Shortage penalty..* Figure 2B gives mean total cost per pair (holding, order-
434 ing and shortage cost) relative to SAP-derived parameters in the hard-limit
435 replay for penalties of one and five times the unit cost. Costs are dominated
436 by the shortage term, and the differences among the non-tuned arms (0.1–
437 3.7%) should be read with caution. The simulation-tuned policy is by far
438 the cheapest (1,094,043 at one times the unit cost, about 33% below SAP-
439 derived), because it was selected to meet the service target on the pre-cutoff
440 history. Among the other arms SAP-derived parameters are the cheapest,
441 as they carry more inventory and fewer stockouts. LLM run 1 then costs
442 2.4% more than SAP-derived at both multiples, is about as expensive as the
443 Croston/SBA OR (1,673,003 versus 1,674,925 at one times), and is cheaper
444 than the SLT-informed (1.2%), ERP-floor, empirical-quantile and history-
445 only (0.8%) ORs. The LLM’s cost advantage over ERP-derived parameters
446 is therefore specific to the case in which service shortfalls are not charged.

447 7. Managerial Insights

- 448 1. **Position the tool as a reviewed decision support system, not**
449 **an autonomous engine.** Single generations differ at the material level
450 (Section 6.7). The gate-simulate-review structure yields a justified can-
451 didate for every material and leaves the decision to the planner; whether
452 this helps planners is untested.
- 453 2. **Benchmark against data-driven baselines.** ERP-derived parameters
454 and ERP-informed ORs are weak benchmarks (safety stocks are often
455 stale or zero; safety lead times inflate buffers), so comparing only with
456 ERP overestimates the benefit: the LLM cost 6%, 8% and 18% more than
457 history-only, Croston/SBA and empirical-quantile baselines.
- 458 3. **Expect plant-level differences.** Against ERP-derived parameters the
459 LLM was cheaper in Plants A and B (14–23%) and on 33 lumpy materials
460 (about 25%, mixed service) but not in Plant C; we did not test whether
461 escalation rules would catch such cases.
- 462 4. **Verify storage along the trajectory.** The static rule $SS + Q \leq \text{limit}$
463 cannot protect a warehouse: 53–66% of pairs breach the limit in simula-
464 tion under each of the three arms checked. Enforce a cap in the order-
465 release logic or treat the limit as a design requirement.
- 466 5. **Price shortages explicitly.** Without a shortage penalty the ranking
467 rewards under-stocking; with a penalty of one unit cost the LLM is about
468 2.4% more expensive than SAP-derived. Set service targets and penalties
469 before reading cost comparisons.
- 470 6. **Freeze and log the configuration.** Treat the system prompt as a

471 controlled file, log every call (input hash, time, session, raw response)
472 and re-run before changing the prompt.

473 **7. Cost and latency.** Each generation took several minutes, whereas OR
474 baselines take milliseconds; periodic regeneration would support adapt-
475 able planning [83].

476 **8. Discussion, Limitations and Research Perspectives**

477 *What the evidence shows..* A prompt-configured LLM generates feasible, re-
478 viewable (r, Q) policies for every material; in simulation they cost less than
479 ERP-derived parameters but more than data-driven OR policies. InventOR
480 is therefore not presented as outperforming optimization: its role is a re-
481 viewed candidate generator, and whether that helps planners is a hypothesis
482 the simulation cannot test. Against the ERP-informed arms alone the LLM
483 would look far better (-11.4% against SAP-derived, -30.5% against the
484 SLT-informed OR).

485 *Design principles for LLM-based decision support..* We derive six principles
486 from this application and offer them as hypotheses for design-science work
487 [7]. **DP1, bound the generator:** the LLM is a candidate generator whose
488 prompt fixes the method, constraints and output schema, with no authority
489 to act. **DP2, verify deterministically and along time:** check outputs
490 with code and feasibility along the trajectory, because point checks can pass
491 while constraints are violated over time (Section 6.5). **DP3, show counter-**
492 **factual evidence:** give the reviewer simulated comparisons with credible
493 references. **DP4, control information in evaluation:** compare against
494 baselines with and without the same information and withhold incumbent
495 inputs in an ablation (Sections 6.3, 6.4). **DP5, freeze and log:** treat the

496 prompt as a controlled file, log every call and repeat runs (Section 6.7). **DP6**,
497 **escalate by rule:** define in advance when proposals are trusted, escalated
498 or rejected, to limit complacency [13] (Section 4.3).

499 *Limitations.* (i) The data come from one manufacturer and three plants. (ii)
500 The model identifier and inference settings are operator-declared. (iii) No
501 planner outcomes were measured: the benefits of reviewability, auditability
502 and planner support and the rules of Section 4.3 are design intentions that
503 a field study must test. (iv) The simulation uses deterministic arrival at the
504 rounded mean lead time, ignores observed receipts, corrections and scrap,
505 starts with an empty pipeline and does not simulate lead-time variability,
506 which may bias comparisons that depend on safety lead time; we did not
507 report service-target attainment. (v) The primary results use a zero shortage
508 penalty; with a positive penalty the ranking changes. Costs are dominated
509 by a few high-demand materials and the 346 pairs are not independent,
510 so the bootstrap intervals understate uncertainty. (vi) The (r, Q) family
511 is single-item and single-echelon. (vii) Individual policies vary across runs,
512 and Plant C required calendar reconstruction. (viii) We included no deep
513 reinforcement-learning or neural forecasting baseline, and the ERP-derived
514 arm (ERP safety stock with our own order quantity) is not an observed
515 record of practice.

516 *Research perspectives.*

517 1. *Ensembling and hybrids.* Aggregating several generations may reduce
518 individual-policy noise [42], and combining the best data-driven OR policy
519 with LLM review or regime-specific choice [84] could capture the OR cost
520 level and the LLM’s flags.

521 2. *Comparison with trained methods.* Benchmarks against deep
 522 reinforcement-learning and other trained policies [53, 46, 47] on
 523 identical data and information, and richer policy families such as (s, S)
 524 [58, 59], would clarify the trade-off between setup effort and policy
 525 quality.

526 3. *Planner-in-the-loop field trials.* Prospective studies should compare plan-
 527 ners with and without the aid on matched materials, measuring planner
 528 time, override rates and realized service and inventory, and should exam-
 529 ine adoption drivers [28, 27, 5].

530 9. Conclusion

531 We defined InventOR as a human-reviewed decision support system for a
 532 recurring production-planning decision, the setting of replenishment param-
 533 eters per plant-material pair, and evaluated it against six design requirements
 534 on a three-plant industrial dataset. A prompt-configured LLM generated
 535 feasible, justified (r, Q) policies for every material without training, and a
 536 deterministic gate, a hard-limit simulation and explicit review rules made
 537 them checkable. The proposals cost about 11% less than ERP-derived pa-
 538 rameters at slightly lower service, but the LLM costs about 6%, 8% and 18%
 539 more than history-only, Croston/SBA and empirical-quantile policies; when
 540 lost sales are priced at one unit cost or more, ERP-derived parameters are
 541 the cheapest of the non-tuned arms and a simulation-tuned policy is cheapest
 542 overall. Storage limits that hold in a static check are exceeded along simu-
 543 lated trajectories for every method, and the LLM loses to ERP at one plant.
 544 The evidence supports InventOR as a checkable candidate generator inside
 545 a reviewed process, not as an optimizer that beats data-driven baselines; the

design principles of Section 8 are hypotheses, and field evidence on planner behaviour and realized outcomes is the central open step.

Appendix A. Formal Definitions

Reference calculation. These are the standard textbook calculations [67, 78]. For pair i , let $\bar{d}_i$ and $\sigma_{D,i}$ be the mean and standard deviation of pre-cutoff working-day demand, L_i and $\sigma_{L,i}$ the mean and standard deviation of the lead time in working days, K_i the ordering cost, h_i the annual holding cost per unit (unit cost times holding rate), z_i the standard normal quantile of the service target, and 250 the working days per year. The order quantity is $Q_i = \max\{\text{MOQ}_i, \lceil \sqrt{2K_i \cdot 250\bar{d}_i/h_i} \rceil\}$. The lead-time-demand standard deviation is $\sigma_{LTD,i} = \sqrt{L_i\sigma_{D,i}^2 + \bar{d}_i^2\sigma_{L,i}^2}$, the safety stock $SS_i = \lceil z_i\sigma_{LTD,i} \rceil$ and the reorder point $R_i = \lceil \bar{d}_iL_i^{\text{eff}} + SS_i \rceil$, where L_i^{eff} is the mean lead time for the history-only OR and the mean lead time plus the ERP safety lead time for the SLT-informed OR and the ERP-floor OR. The SAP-derived arm uses the ERP safety stock and the same Q_i . All policies are projected to $SS_i + Q_i \leq \text{max_storage_units}_i$ when $\text{MOQ}_i + SS_i$ fits.

Simulation. On each post-cutoff working day, demand is served from on-hand stock, unmet demand is lost, and an order of size Q_i is placed when the inventory position (on-hand plus on order) is at or below R_i and arrives after the rounded mean lead time in subsequent working observations. Observed receipts, corrections and scrap are ignored. The total cost is $TC_i = h_i\bar{I}_iT/250 + K_iN_i$, where $\bar{I}_i$ is the mean on-hand stock over the T simulated working days and N_i the number of orders. Fill rate is served divided by demanded units. In the hard-limit replay an order is truncated to $\text{max_storage_units}_i$ minus the inventory position, and no order is placed

571 if the truncated quantity is zero.

572 **Author Contributions**

573 Conceptualization, Methodology, Software, Formal analysis, Validation,
574 Visualization, Writing – original draft, Writing – review and editing: Aris
575 Dressino.

576 **Funding**

577 This research did not receive any specific grant from funding agencies in
578 the public, commercial, or not-for-profit sectors.

579 **Declaration of Competing Interest**

580 The author declares no known competing financial interests or personal
581 relationships that could have appeared to influence the work reported in this
582 paper.

583 **Data and Code Availability**

584 The raw operational data and material-level artifacts cannot be publicly
585 released because they contain commercially sensitive material, supplier
586 and planning information. A confidentiality-screened release package
587 containing analysis code, the anonymized variable schema, aggregate
588 evaluation results, manuscript sources and prompt files is available at
589 github.com/DressPD/inventor_public. A versioned Zenodo archive of
590 the release is available at <https://doi.org/10.5281/zenodo.22262052>.
591 Supplementary File S1 contains the system prompt used for all runs; the
592 prompt file, the runner code and, for each call, a hash of the input, the

593 UTC time, the session identifier and the raw response are in the release.
594 The model identifier and inference settings are not exposed by the platform
595 and are operator-declared (Anthropic Claude Sonnet 4.6).

596 **Acknowledgments**

597 The author thanks the industrial collaborators and material-control ex-
598 perts who contributed domain knowledge and access to anonymized opera-
599 tional data, and the enterprise data and IT teams for supporting data access
600 and prototype execution. Any errors or omissions remain the responsibility
601 of the author.

602 **Declaration of Generative AI and AI-Assisted Technologies in the** 603 **Writing Process**

604 During manuscript preparation and revision, the author used OpenAI
605 GPT-5.6 and Anthropic Claude (Sonnet 5.5) models through the OpenCode
606 coding assistant for language editing, restructuring, consistency checks, ci-
607 tation verification against Crossref metadata, LaTeX proofing, and writing
608 and executing the analysis scripts that produce the evidence tables. These
609 assistants are distinct from the Claude Sonnet 4.6 deployment evaluated in
610 the study. The author specified the analyses, reviewed and edited all content
611 and results, and takes full responsibility for the content. No generative AI
612 created or manipulated image content.

613 **References**

614 [1] I. Jackson, D. Ivanov, A. Dolgui, J. Namdar, Generative artificial intelligence
615 in supply chain and operations management: A capability-based framework

- 616 for analysis and implementation, *International Journal of Production Research*
617 62 (17) (2024) 6120–6145. doi:10.1080/00207543.2024.2309309.
- 618 [2] S. Fosso Wamba, C. Guthrie, M. M. Queiroz, S. Minner, ChatGPT and gener-
619 ative artificial intelligence: An exploratory study of key benefits and challenges
620 in operations and supply chain management, *International Journal of Produc-*
621 *tion Research* 62 (16) (2024) 5676–5696. doi:10.1080/00207543.2023.2294116.
- 622 [3] Z. Song, Y. Xie, L. Yang, Y. Zhao, Large language models in supply
623 chain management: A systematic literature review and application frame-
624 work, *International Journal of Production Research* 64 (16) (2026) 7144–7184.
625 doi:10.1080/00207543.2026.2641103.
- 626 [4] S. Panja, T. Zheng, C. H. Glock, Potentials of generative AI in logistics and
627 supply chain management: A guiding review on applications and future re-
628 search directions, *International Journal of Production Research* Advance online
629 publication (2026). doi:10.1080/00207543.2026.2680586.
- 630 [5] N. Trifone, F. Dallari, A. Brintrup, Drivers and constraints of agen-
631 tic AI in supply chain management: An exploratory study, *Interna-*
632 *tional Journal of Production Research* Advance online publication (2026).
633 doi:10.1080/00207543.2026.2669878.
- 634 [6] Z. Ji, N. Lee, R. Frieske, T. Yu, D. Su, Y. Xu, E. Ishii, Y. J. Bang, A. Madotto,
635 P. Fung, Survey of hallucination in natural language generation, *ACM Com-*
636 *puting Surveys* 55 (12) (2023) 1–38. doi:10.1145/3571730.
- 637 [7] A. R. Hevner, S. T. March, J. Park, S. Ram, Design science in information
638 systems research, *MIS Quarterly* 28 (1) (2004) 75–105. doi:10.2307/25148625.
- 639 [8] K. Peffers, T. Tuunanen, M. A. Rothenberger, S. Chatterjee, A design sci-
640 ence research methodology for information systems research, *Journal of Man-*
641 *agement Information Systems* 24 (3) (2007) 45–77. doi:10.2753/MIS0742-
642 1222240302.

- 643 [9] G. A. Gorry, M. S. Scott Morton, A framework for management information
644 systems, *Sloan Management Review* 13 (1) (1971) 55–70.
- 645 [10] R. H. Sprague, A framework for the development of decision support systems,
646 *MIS Quarterly* 4 (4) (1980) 1–26. doi:10.2307/248957.
- 647 [11] J. P. Shim, M. Warkentin, J. F. Courtney, D. J. Power, R. Sharda, C. Carlsson,
648 Past, present, and future of decision support technology, *Decision Support*
649 *Systems* 33 (2) (2002) 111–126. doi:10.1016/S0167-9236(01)00139-7.
- 650 [12] R. Parasuraman, T. B. Sheridan, C. D. Wickens, A model for types and lev-
651 els of human interaction with automation, *IEEE Transactions on Systems,*
652 *Man, and Cybernetics – Part A: Systems and Humans* 30 (3) (2000) 286–297.
653 doi:10.1109/3468.844354.
- 654 [13] R. Parasuraman, D. H. Manzey, Complacency and bias in human use of au-
655 tomation: An attentional integration, *Human Factors* 52 (3) (2010) 381–410.
656 doi:10.1177/0018720810376055.
- 657 [14] S. Amershi, D. Weld, M. Vorvoreanu, A. Fourney, B. Nushi, P. Collisson,
658 J. Suh, S. Iqbal, P. N. Bennett, K. Inkpen, et al., Guidelines for human-AI
659 interaction, in: *Proceedings of the 2019 CHI Conference on Human Factors in*
660 *Computing Systems*, 2019, pp. 1–13. doi:10.1145/3290605.3300233.
- 661 [15] I. D. Raji, A. Smart, R. N. White, M. Mitchell, T. Gebru, B. Hutchinson,
662 J. Smith-Loud, D. Theron, P. Barnes, Closing the AI accountability gap:
663 Defining an end-to-end framework for internal algorithmic auditing, in: *Pro-*
664 *ceedings of the 2020 Conference on Fairness, Accountability, and Transparency*,
665 2020, pp. 33–44. doi:10.1145/3351095.3372873.
- 666 [16] S. Ouyang, J. M. Zhang, M. Harman, M. Wang, An empirical study of the non-
667 determinism of ChatGPT in code generation, *ACM Transactions on Software*
668 *Engineering and Methodology* 34 (2) (2025) 1–28. doi:10.1145/3697010.

- 669 [17] T. Schelhorn, U. Gnewuch, A. Maedche, Effective use of open data portals for
670 social good: Design principles for explainable LLM-based open data assistants,
671 Decision Support Systems 207 (2026) 114698. doi:10.1016/j.dss.2026.114698.
- 672 [18] D. Prithvirajan, M. Mathirajan, An AI-enabled decision support sys-
673 tem for evidence-based last-mile delivery authentication: Design and
674 pilot implementation, Decision Support Systems 211 (2026) 114760.
675 doi:10.1016/j.dss.2026.114760.
- 676 [19] X. Ning, Y. Lu, W. Li, S. Gupta, How transparency affects algorithmic advice
677 utilization: The mediating roles of trusting beliefs, Decision Support Systems
678 183 (2024) 114273. doi:10.1016/j.dss.2024.114273.
- 679 [20] L. Soleimanof, D. Neufeld, AI advice and human metacognition, Decision Sup-
680 port Systems 210 (2026) 114757. doi:10.1016/j.dss.2026.114757.
- 681 [21] B. Chae, C. Sheu, E. Park, The value of data, machine learning, and
682 deep learning in restaurant demand forecasting: Insights and lessons learned
683 from a large restaurant chain, Decision Support Systems 184 (2024) 114291.
684 doi:10.1016/j.dss.2024.114291.
- 685 [22] A. Engin, L. A. Franco, E. A. J. A. Rouwette, How groups manage con-
686 flict when using model-driven decision support: An epistemic motivation lens,
687 Omega 129 (2024) 103142. doi:10.1016/j.omega.2024.103142.
- 688 [23] X. Wang, V. S. Rodrigues, E. Demir, J. Sarkis, Algorithm aversion during
689 disruptions: The case of safety stock, International Journal of Production
690 Economics 278 (2024) 109442. doi:10.1016/j.ijpe.2024.109442.
- 691 [24] S. Fosso Wamba, M. M. Queiroz, C. J. Chiappetta Jabbour, C. Shi, Are both
692 generative AI and ChatGPT game changers for 21st-century operations and
693 supply chain excellence?, International Journal of Production Economics 265
694 (2023) 109015. doi:10.1016/j.ijpe.2023.109015.

- 695 [25] L. Li, Y. Liu, Y. Jin, T. C. E. Cheng, Q. Zhang, Generative AI-
696 enabled supply chain management: The critical role of coordination and dy-
697 namism, *International Journal of Production Economics* 277 (2024) 109388.
698 doi:10.1016/j.ijpe.2024.109388.
- 699 [26] T. Boone, B. Fahimnia, R. Ganeshan, D. M. Herold, N. R. Sanders, Gener-
700 ative AI: Opportunities, challenges, and research directions for supply chain
701 resilience, *Transportation Research Part E: Logistics and Transportation Re-*
702 *view* 199 (2025) 104135. doi:10.1016/j.tre.2025.104135.
- 703 [27] V. G. Cannas, M. P. Ciano, M. Saltalamacchia, R. Secchi, Artificial intelli-
704 gence in supply chain and operations management: A multiple case study re-
705 search, *International Journal of Production Research* 62 (9) (2024) 3333–3360.
706 doi:10.1080/00207543.2023.2232050.
- 707 [28] V. Venkatesh, R. Raman, F. Cruz-Jesus, AI and emerging technol-
708 ogy adoption: A research agenda for operations management, *Inter-*
709 *national Journal of Production Research* 62 (15) (2024) 5367–5377.
710 doi:10.1080/00207543.2023.2192309.
- 711 [29] M. C. Cohen, T. Dai, G. Perakis, N. Agrawal, G. Allon, R. N. Boute, G. P.
712 Cachon, Z. Chen, M. Cohen, R. Cristian, et al., OM Forum—Supply Chain
713 Management in the AI Era: A vision statement from the operations manage-
714 ment community, *Manufacturing & Service Operations Management* 28 (3)
715 (2026) 687–705. doi:10.1287/msom.2025.1065.
- 716 [30] F. Olan, K. Spanaki, W. Ahmed, G. Zhao, Enabling explainable artificial intel-
717 ligence capabilities in supply chain decision support making, *Production Plan-*
718 *ning & Control* 36 (6) (2025) 808–819. doi:10.1080/09537287.2024.2313514.
- 719 [31] P. Amarasinghe, S. Nguyen, Y. Sun, S. Arisian, D. Alahakoon, Business op-
720 timization for digital manufacturing: A fine-tuned large language model ap-
721 proach, *International Journal of Production Economics* 295 (2026) 109934.
722 doi:10.1016/j.ijpe.2026.109934.

- [32] M. Liu, Y. Bai, M. Qi, Z.-J. M. Shen, Inventory management with transformer: Automated decision making for order timing and quantity, *Service Science* 18 (3) (2026) 196–211. doi:10.1287/serv.2024.0236.
- [33] Y. Chen, S. N. Kirshner, A. Ovchinnikov, M. Andiappan, T. Jenkin, A manager and an AI walk into a bar: Does ChatGPT make biased decisions like we do?, *Manufacturing & Service Operations Management* 27 (2) (2025) 354–368. doi:10.1287/msom.2023.0279.
- [34] B. Li, K. Mellou, B. Zhang, J. Pathuri, I. Menache, Large language models for supply chain optimization, *arXiv preprint arXiv:2307.03875* (2023). doi:10.48550/arXiv.2307.03875.
- [35] C. Huang, Z. Tang, S. Hu, R. Jiang, X. Zheng, D. Ge, B. Wang, Z. Wang, ORLM: A customizable framework in training large models for automated optimization modeling, *Operations Research* 73 (6) (2025) 2986–3009. doi:10.1287/opre.2024.1233.
- [36] A. Ahmaditeshnizi, W. Gao, M. Udell, OptiMUS: Scalable optimization modeling with (MI)LP solvers and large language models, in: *Proceedings of the 41st International Conference on Machine Learning*, Vol. 235 of *Proceedings of Machine Learning Research*, PMLR, 2024, pp. 577–596. URL <https://proceedings.mlr.press/v235/ahmaditeshnizi24a.html>
- [37] B. Zhang, P. Luo, G. Yang, B.-H. Soong, C. Yuen, OR-LLM-Agent: Automating modeling and solving of operations research optimization problems with reasoning LLM, *arXiv preprint arXiv:2503.10009* (2025). doi:10.48550/arXiv.2503.10009.
- [38] Y. Quan, Z. Liu, InvAgent: A large language model based multi-agent system for inventory management in supply chains, *arXiv preprint arXiv:2407.11384* (2024). doi:10.48550/arXiv.2407.11384.

- [39] R. Jannelli, M. Schöpf, L. Bickel, T. H. Netland, A. Brintrup, Agentic LLMs in the supply chain: Towards autonomous multi-agent consensus-seeking, *International Journal of Production Research* Advance online publication (2025). doi:10.1080/00207543.2025.2604311.
- [40] K.-Y. Lin, S.-Y. Wu, K. Matsuno, Generative AI-driven resilience in supply chain management: UNISONE framework for disruption modelling and capacity optimisation, *International Journal of Production Research* Advance online publication (2025). doi:10.1080/00207543.2025.2574412.
- [41] S. Hamdan, N. Brahimi, Evaluating large language models for complex supply chain optimisation: A focus on routing problems, *International Journal of Production Research* Advance online publication (2026). doi:10.1080/00207543.2026.2737946.
- [42] M. Abbaszadeh Nakhost, A. Bloemer, S. Minner, Full automation or just a copilot? a study of large language models on spare parts demand forecasting with OpenAI o1-mini, *International Journal of Production Research* Advance online publication (2025). doi:10.1080/00207543.2025.2591225.
- [43] K. Ren, H. Bidkhori, Z.-J. M. Shen, Data-driven inventory policy: Learning from sequentially observed non-stationary data, *Omega* 123 (2024) 102942. doi:10.1016/j.omega.2023.102942.
- [44] Y. R. Sagaert, N. Kourentzes, Inventory management with leading indicator augmented hierarchical forecasts, *Omega* 136 (2025) 103335. doi:10.1016/j.omega.2025.103335.
- [45] S. Gulecyuz, B. O’Sullivan, S. Armagan Tarim, A heuristic method for perishable inventory management under non-stationary demand, *Omega* 133 (2025) 103267. doi:10.1016/j.omega.2024.103267.
- [46] M. Kolyaei, L. Zhang, M. Blom, Inventory replenishment and fulfilment decisions for an omnichannel retailer: A reinforcement learning-based method,

- International Journal of Production Research 63 (24) (2025) 9571–9592.
doi:10.1080/00207543.2025.2520596.
- [47] F. Stranieri, F. Stella, C. Kouki, Performance of deep reinforcement learning algorithms in two-echelon inventory control systems, International Journal of Production Research 62 (17) (2024) 6211–6226.
doi:10.1080/00207543.2024.2311180.
- [48] S. Ma, Omni-channel retail inventory management with deep reinforcement learning, Omega 144 (2026) 103586. doi:10.1016/j.omega.2026.103586.
- [49] H. Wang, J. Tao, T. Peng, A. Brintrup, E. E. Kosasih, Y. Lu, R. Tang, L. Hu, Dynamic inventory replenishment strategy for aerospace manufacturing supply chain: Combining reinforcement learning and multi-agent simulation, International Journal of Production Research 60 (13) (2022) 4117–4136.
doi:10.1080/00207543.2021.2020927.
- [50] S. Ma, Integrated demand forecasting and inventory replenishment in omni-channel retail with deep learning, International Journal of Production Research Advance online publication (2026). doi:10.1080/00207543.2026.2738209.
- [51] Q. Wang, Y. Chen, Y. Mi, C. She, P. Zhao, C. Chen, Contextual distributionally robust optimisation for inventory replenishment optimisation under uncertain demand, International Journal of Production Research 64 (19) (2026) 8393–8411. doi:10.1080/00207543.2026.2653131.
- [52] X. Zhang, W. Dai, X. Cai, Inventory replenishment decisions with uncertain price and demand, International Journal of Production Research 62 (7) (2024) 2663–2682. doi:10.1080/00207543.2023.2223701.
- [53] M. Gabellini, A. Regattieri, M. Bortolini, M. Ronchi, A novel data-driven approach for integrated inventory replenishment, safety stock and lead time optimisation, International Journal of Production Research Advance online publication (2026). doi:10.1080/00207543.2026.2696440.

- [54] G. Schett, C. Ehrig, B. Schwendinger, U. Neumann, N. Beck, F. Stadler, F. Ansari, Integrated demand forecasting and reinforcement learning for order point optimization in inventory planning, *International Journal of Production Economics* 299 (2026) 110111. doi:10.1016/j.ijpe.2026.110111.
- [55] J. D. Croston, Forecasting and stock control for intermittent demands, *Operational Research Quarterly* 23 (3) (1972) 289–303. doi:10.1057/jors.1972.50.
- [56] A. A. Syntetos, J. E. Boylan, The accuracy of intermittent demand estimates, *International Journal of Forecasting* 21 (2) (2005) 303–314. doi:10.1016/j.ijforecast.2004.10.001.
- [57] T. R. Willemain, C. N. Smart, H. F. Schwarz, A new approach to forecasting intermittent demand for service parts inventories, *International Journal of Forecasting* 20 (3) (2004) 375–387. doi:10.1016/S0169-2070(03)00013-X.
- [58] D. Prak, R. Teunter, M. Z. Babai, J. E. Boylan, A. Syntetos, Robust compound Poisson parameter estimation for inventory control, *Omega* 104 (2021) 102481. doi:10.1016/j.omega.2021.102481.
- [59] M. Z. Babai, A. A. Syntetos, R. H. Teunter, Fifty years of inventory research from a forecasting perspective, *European Journal of Operational Research* 331 (1) (2026) 1–20. doi:10.1016/j.ejor.2025.07.003.
- [60] K. Sanguri, S. Patra, K. Nikolopoulos, S. Punia, Intermittent demand, inventory obsolescence, and temporal aggregation forecasts, *International Journal of Production Research* 62 (5) (2024) 1663–1685. doi:10.1080/00207543.2023.2199435.
- [61] L. Li, Y. Kang, F. Petropoulos, F. Li, Feature-based intermittent demand forecast combinations: Accuracy and inventory implications, *International Journal of Production Research* 61 (22) (2023) 7557–7572. doi:10.1080/00207543.2022.2153941.

- [62] R. L. Goldsmith, A.-L. Sachs, Spare part demand forecasting in every phase: A data pooling approach to the Bass life cycle model, *International Journal of Production Research* Advance online publication (2025). doi:10.1080/00207543.2025.2607643.
- [63] V. Babaveisi, E. Teimoury, M. R. Gholamian, B. Rostami-Tabar, Integrated demand forecasting and planning model for repairable spare part: An empirical investigation, *International Journal of Production Research* 61 (20) (2023) 6791–6807. doi:10.1080/00207543.2022.2137596.
- [64] M. Caserta, L. D’Angelo, Optimising two-echelon spare parts inventory: A case study from Amazon’s operations, *International Journal of Production Research* 64 (19) (2026) 8425–8443. doi:10.1080/00207543.2026.2655026.
- [65] G. Hadley, T. M. Whitin, *Analysis of Inventory Systems*, Prentice-Hall, Englewood Cliffs, NJ, 1963.
- [66] E. L. Porteus, *Foundations of Stochastic Inventory Theory*, Stanford University Press, Stanford, CA, 2002. doi:10.1515/9781503619883.
- [67] E. A. Silver, D. F. Pyke, D. J. Thomas, *Inventory and Production Management in Supply Chains*, 4th Edition, CRC Press, Boca Raton, FL, 2017. doi:10.1201/9781315374406.
- [68] P. Sahoo, A. K. Singh, S. Saha, V. Jain, S. Mondal, A. Chadha, A systematic survey of prompt engineering in large language models: Techniques and applications, *arXiv preprint arXiv:2402.07927* (2024). doi:10.48550/arXiv.2402.07927.
- [69] L. J. Tashman, Out-of-sample tests of forecasting accuracy: an analysis and review, *International Journal of Forecasting* 16 (4) (2000) 437–450. doi:10.1016/S0169-2070(00)00065-0.

- [70] H. Hewamalage, K. Ackermann, C. Bergmeir, Forecast evaluation for data scientists: common pitfalls and best practices, *Data Mining and Knowledge Discovery* 37 (2) (2023) 788–832. doi:10.1007/s10618-022-00894-5.
- [71] J. Venable, J. Pries-Heje, R. Baskerville, FEDS: A framework for evaluation in design science research, *European Journal of Information Systems* 25 (1) (2016) 77–89. doi:10.1057/ejis.2014.36.
- [72] S. Watts, G. Shankaranarayanan, A. Even, Data quality assessment in context: A cognitive perspective, *Decision Support Systems* 48 (1) (2009) 202–211. doi:10.1016/j.dss.2009.07.012.
- [73] N. DeHoratius, A. Raman, Inventory record inaccuracy: An empirical analysis, *Management Science* 54 (4) (2008) 627–641. doi:10.1287/mnsc.1070.0789.
- [74] Y. Kang, S. B. Gershwin, Information inaccuracy in inventory systems: stock loss and stockout, *IIE transactions* 37 (9) (2005) 843–859. doi:10.1080/07408170590969861.
- [75] A. J. Mersereau, Information-sensitive replenishment when inventory records are inaccurate, *Production and Operations Management* 22 (4) (2013) 792–810. doi:10.1111/j.1937-5956.2011.01305.x.
- [76] T. J. van Kampen, D. P. van Donk, D.-J. van der Zee, Safety stock or safety lead time: coping with unreliability in demand and supply, *International Journal of Production Research* 48 (24) (2010) 7463–7481. doi:10.1080/00207540903348346.
- [77] M. Thürer, N. O. Fernandes, S. Haeussler, M. Stevenson, Dynamic planned lead times in production planning and control systems: Does the lead time syndrome matter?, *International Journal of Production Research* 61 (4) (2023) 1268–1282. doi:10.1080/00207543.2022.2034193.
- [78] S. Axsäter, *Inventory Control*, 3rd Edition, Vol. 225 of International Series in Operations Research & Management Science, Springer, Cham, 2015.

- [79] P. H. Zipkin, Foundations of Inventory Management, McGraw-Hill, New York, 2000.
- [80] M. Bijvank, I. F. A. Vis, Lost-sales inventory theory: A review, *European Journal of Operational Research* 215 (1) (2011) 1–13. doi:10.1016/j.ejor.2011.02.004.
- [81] E. T. Grasso, B. W. Taylor, A simulation-based experimental investigation of supply/timing uncertainty in MRP systems, *International Journal of Production Research* 22 (3) (1984) 485–497. doi:10.1080/00207548408942468.
- [82] H. S. Heese, Minimum order quantity requirements in the presence of supply disruption risk and retail competition, *International Journal of Production Research* 63 (16) (2025) 6056–6070. doi:10.1080/00207543.2025.2468391.
- [83] S. Sengupta, P. Jonsson, H. Dreyer, R. Kaipia, T. Y. Choi, Uncertainty regulation and adaptable supply chain planning, *International Journal of Operations & Production Management* 45 (11) (2025) 1861–1883. doi:10.1108/ijopm-07-2025-0645.
- [84] S. Wasserkrug, L. Boussioux, D. Den Hertog, F. Mirzazadeh, Ş. İ. Birbil, J. Kurtz, D. Maragno, Enhancing decision making through the integration of large language models and operations research optimization, in: *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 39, 2025, pp. 28643–28650. doi:10.1609/aaai.v39i27.35090.